



GPU Teaching Kit
Accelerated Computing



西南石油大学 计算机科学学院

SCHOOL OF COMPUTER SCIENCE, SOUTHWEST PETROLEUM UNIVERSITY



Module 10 – Parallel patterns: Sparse Matrix Computation

Objective

- Understanding the wealth of work in sparse matrix storage formats and their corresponding parallel algorithms.
 - Addressing compaction and regularization challenges.
 - Stored format: **CSR**, ELL, COO, Hybrid approach, JDS
 - The parallel SpMV computation performance



GPU Teaching Kit
Accelerated Computing



西南石油大学 计算机科学学院
SCHOOL OF COMPUTER SCIENCE, SOUTHWEST PETROLEUM UNIVERSITY



Module 10 – Parallel patterns: Sparse Matrix Computation

10.1 Background

Objective

- **Understanding what a sparse matrix is along with its storage format and implementing SpMV based on the CSR format.**
 - **Compressed Sparse Row (CSR) storage format**
 - **Implementing SpMV based on the CSR format.**

What we need to know before learning

- **A sparse matrix** is a matrix where **the majority of the elements are zeros**.
- **Sparse matrices** appear in many scientific, engineering, and financial modeling problems.
- Each row of the matrix represents **an equation of a linear system**.
- In various scientific and engineering problems, a large number of equations involve only a small number of variables.

Compressed Sparse Row (CSR) storage format

Sparse Matrix: A

Row 0	3	0	1	0
Row 1	0	0	0	0
Row 2	0	2	4	1
Row 3	1	0	0	1

		Row 0	Row 2	Row 3	
Nonzero values	data[7]	{ 3, 1,	2, 4, 1,	1, 1 }	
Column indices	col_index[7]	{ 0, 2,	1, 2, 3,	0, 3 }	
Row Pointers	row_ptr[5]	{ 0, 2, 2, 5, 7 }			
		Row 0	Row 1	Row 2	Row 4

Example of Compressed Sparse Row (CSR) format.

Why We Need Sparse Matrix–Vector Multiplication (SpMV)

- Solving a linear system of N equations of N variables in the form:

$$A * X + Y = 0,$$

- where A is an $N \times N$ matrix, X is a vector of N variables, and Y is a vector of N constant values. **The objective is to solve for the X variable** that will satisfy all the equations.
- An intuitive (直观) approach is to inverse the matrix such that:

$$X = A^{-1} * (-Y)$$

Gaussian elimination


高斯消元法

- Instead, linear systems of equations represented in sparse matrices can be better solved with an iterative approach:

$$A * X + Y \rightarrow 0$$

Sparse Matrix–Vector Multiplication (SpMV)

$$A * X + Y \rightarrow Y$$



An example of matrix–vector multiplication and accumulation.

A Sequential SpMV/CSR

```
for (int row = 0; row < num_rows; row++) {
```

```
    float dot = 0;
```

```
    int row_start = row_ptr[row];
```

```
    int row_end = row_ptr[row+1];
```

```
    for (int elem = row_start; elem < row_end; elem++) {
```

```
        dot += data[elem] * x[col_index[elem]];
```

```
    }
```

```
    y[row] += dot;
```

```
}
```

Sparse Matrix: A

Row 0

3	0	1	0
---	---	---	---

Row 1


0	0	0	0
---	---	---	---

Row 2

0	2	4	1
---	---	---	---

Row 3

1	0	0	1
---	---	---	---



X

1
2
3
4

\hat{X}

6
0
20
5



GPU Teaching Kit
Accelerated Computing



西南石油大学 计算机科学学院
SCHOOL OF COMPUTER SCIENCE, SOUTHWEST PETROLEUM UNIVERSITY



Module 10 – Parallel patterns: Sparse Matrix Computation


10.2 Parallel SPMV using CSR

Objective

- **Understanding Parallel SPMV using CSR.**
 - Mapping threads to rows.
 - Performance analysis

Mapping Threads to Rows

Thread 0	3	0	1	0
Thread 1	0	0	0	0
Thread 2	0	2	4	1
Thread 3	1	0	0	1



A parallel SpMV/CSR kernel

```
__global__ void SpMV_CSR(int num_rows, float *data, int *col_index,
    int *row_ptr, float *x, float *y) {

    int row = blockIdx.x * blockDim.x + threadIdx.x;

    if (row < num_rows) {

        float dot = 0;

        int row_start = row_ptr[row];

        int row_end = row_ptr[row+1];

        for (int elem = row_start; elem < row_end; elem++) {

            dot += data[elem] * x[col_index[elem]];



        }

        y[row] += dot;

    }
}
```

Shortcomings of SpMV/CSR Kernel

- The parallel SpMV/CSR kernel has two major shortcomings
 - The kernel does not make coalesced memory accesses.



Shortcomings of SpMV/CSR Kernel(cont.)

- The parallel SpMV/CSR kernel has two major shortcomings
 - its potential to incur significant control flow divergence in all warps.

Thread 0	3	0	1	0	2	No. of nonzero elements
Thread 1	0	0	0	0	0	
Thread 2	0	2	4	1	3	
Thread 3	1	0	0	1	2	

Adjacent rows can have varying numbers of nonzero elements



GPU Teaching Kit
Accelerated Computing



西南石油大学 计算机科学学院
SCHOOL OF COMPUTER SCIENCE, SOUTHWEST PETROLEUM UNIVERSITY



Module 10 – Parallel patterns: sparse matrix computation

10.3 Padding and Transposition

Objective


- Understanding how to use ELL to address the problems.
 - How to address the problems of noncoalesced memory accesses and control divergence.
 - ELL storage format.

How to solve the two problems

- By applying data padding and transpose on sparse matrix data.
- These ideas were used in the **ELL** storage format, whose name came from the sparse matrix package in **ELLPACK**, a package for solving elliptic boundary value problems

ELL : PADDING AND TRANSPOSITION

- Determine the row with the maximum number of non-zero elements.
- Add pseudo (zero) elements to all other rows after non-zero elements, so that they have **the same length as the maximum row**.
- Row priority **transpose** rectangular matrix.



Example of ELL format.

Row 0	3	0	1	0
Row 1	0	0	0	0
Row 2	0	2	4	1
Row 3	1	0	0	1

Step1:CSR with padding

Values	Columns
3	0
1	2
*	*
*	*
2	1
4	2
1	3
*	*
1	0
1	3
*	*

Step2:transpose

Values	Columns
3	0
*	*
2	1
1	2
*	*
4	2
1	3
*	*
1	3
*	*

A	{	data	3	*	2	1	1	*	4	1	*	*	1	*
		col_index	0	*	1	0	2	*	2	3	*	*	3	*

ELL

Note: We no longer need row_ptr array, because the start of the **i-th row** (which is now the **i-th column** in transposed matrix) has been simplified to **data [i]** through **data[num_rows-1]**.

Example of ELL format.


Row 0	3	0	1	0
Row 1	0	0	0	0
Row 2	0	2	4	1
Row 3	1	0	0	1

Step1:CSR with padding

Values			Columns		
3	1	*	0	2	*
*	*	*	*	*	*
2	4	1	1	2	3
1	1	*	0	3	*

Step2:transpose

Values				Columns			
3	*	2	1	0	*	1	0
1	*	4	1	2	*	2	3
*	*	1	*	*	*	3	*



Example of ELL format.


Row 0	3	0	1	0
Row 1	0	0	0	0
Row 2	0	2	4	1
Row 3	1	0	0	1

Step1:CSR with padding

Values	Columns
3	0
1	2
*	*
*	*
2	1
4	2
1	3
*	*
0	3
*	*

Step2:transpose

Values	Columns
3	0
*	*
2	1
1	2
*	*
4	2
1	3
*	*
1	3
*	*



Example of ELL format.


Row 0	3	0	1	0
Row 1	0	0	0	0
Row 2	0	2	4	1
Row 3	1	0	0	1

Step1:CSR with padding

Values	Columns
3	0
1	2
*	*
*	*
*	*
2	1
4	2
1	3
*	*
1	0
1	3
*	*


Step2:transpose

Values	Columns
3	0
*	*
2	1
1	2
*	*
4	2
1	3
*	*
1	0
*	*
*	3
*	*



A parallel SpMV/ELL kernel

Row 0	3	0	1	0
Row 1	0	0	0	0
Row 2	0	2	4	1
Row 3	1	0	0	1



num_rows=4; //the number of rows in the sparse matrix

num_elem=3; //the maximum number of non-zero elements in all rows of the original sparse matrix.

```
__global__ void SpMV_ELL( int num_rows, float *data, int *col_index, int num_elem,
float *x, float *y)
{
    int row = blockIdx.x * blockDim.x + threadIdx.x;
    int col;
    if (row < num_rows) {
        float dot = 0;
        for (int i = 0; i < num_elem; i++) {
            col = col_index[row+i*num_rows];
            dot += data[row+i*num_rows] * x[col];
        }
        y[row] += dot;
    }
}
```


Disadvantages of SpMV/ELL Kernel

- Unfortunately, **SpMV/ELL has a potential drawback**. In cases where one or a few lines have a **significant number of non zero elements**, the ELL format will result in **too many padding elements**.
- This will create a performance bottleneck and storage space issue, which is the problem of uneven thread block load.

How to solve this problem?



GPU Teaching Kit
Accelerated Computing



西南石油大学 计算机科学学院
SCHOOL OF COMPUTER SCIENCE, SOUTHWEST PETROLEUM UNIVERSITY



Module 10 – Parallel patterns: sparse matrix computation

10.4 Using a Hybrid Approach to Regulate Padding

Objective

- **Understanding of how to use a hybrid approach to regulate padding.**
 - The root of the problem with excessive padding in the ELL representation is that one or a small number of rows have an exceedingly large number of nonzero elements.
 - **The Coordinate (COO) format** provides such a mechanism.

COO

Row 0	3	0	1	0
Row 1	0	0	0	0
Row 2	0	2	4	1
Row 3	1	0	0	1

Each non-zero element is stored together with its column index and row index.

			Row 0	Row 2	Row 3
Nonzero values	data[7]	{	3, 1,	2, 4, 1,	1, 1 }
Column indices	col_index[7]	{	0, 2,	1, 2, 3,	0, 3 }
Row indices	row_index[7]	{	0, 0,	2, 2, 2,	3, 3 }

COO

– The characteristics of COO:

- The elements in COO format can be reordered arbitrarily without losing any information.
- You can view any element in the storage and know where it comes from the original sparse matrix.

Nonzero values	data[7]	{ 1 1, 2, 4, 3, 1 1 }
Column indices	col_index[7]	{ 0 2, 1, 2, 0, 3, 3 }
Row indices	row_index[7]	{ 3 0, 2, 2, 0, 2, 3 }

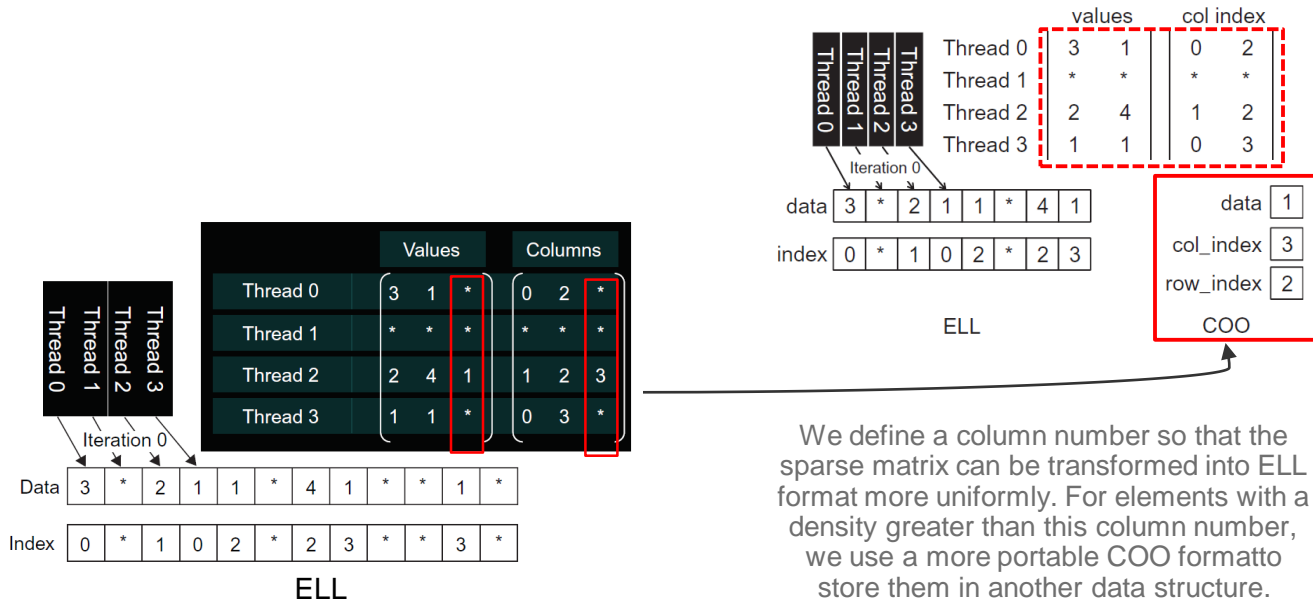
Reordering the COO format.

Limitation: Such reordering would disturb the locality and sequential patterns necessary for the efficient use of memory bandwidth.

Benefit: It can be used to curb the length of rows in the CSR format or the ELL format.

Improve ELL

- We can remove some elements from rows with a very large number of non-zero elements and place them in separate **COO** storage. So we can use SpMV/ELL on other elements.



SpMV kernel using ELL-COO hybrid format

```
//SpMV_Hybrid: Sparse Matrix-Vector Multiplication (SpMV) using a hybrid format
__global__ void SpMV_Hybrid(int num_rows,      // Number of rows in the matrix
    float *data_ell,      // ELL format data array
    int *col_index_ell,   // ELL format column index array
    int num_elem_ell,     // Maximum number of non-zero elements
    float *data_coo,      // COO format data array
    int *row_index_coo,   // COO format row index array
    int *col_index_coo,   // COO format column index array
    int num_elem_coo,     // Number of non-zero elements in COO
    float *x,             // Dense input vector x
    float *y) {           // Output vector y
    // Compute the row index this thread is responsible for
    int row = blockIdx.x * blockDim.x + threadIdx.x;
```

...

SpMV kernel using ELL-COO hybrid format

//SpMV_Hybrid: Sparse Matrix-Vector Multiplication (SpMV) using a hybrid format

...

```
if (row < num_rows) { // Ensure the thread processes a valid row
    float dot = 0;
```

```
// -----
```

data_ell

3	*	2	1	1	*	4	1
---	---	---	---	---	---	---	---

```
// 1. Process using ELL format
```

col_index_ell

0	*	1	0	2	*	2	3
---	---	---	---	---	---	---	---

```
// -----
```

```
for (i = 0; i < num_elem_ell; i++) {
```

```
    // Compute the column index of the i-th element in the current row
```

```
    int col = col_index_ell[i * num_rows + row];
```

```
    if (col != -1) { // Check if the ELL element is valid (-1 indicates no element)
```

```
        // Accumulate the product of the current element and the corresponding
```

```
        // vector x element
```

```
        dot += data_ell[i * num_rows + row] * x[col];
```

```
    }
```

```
}
```


SpMV kernel using ELL-COO hybrid format

//SpMV_Hybrid: Sparse Matrix-Vector Multiplication (SpMV) using a hybrid format

...

// -----

// 2. Process using COO format

// -----

for (i = 0; i < num_elem_coo; i++) {

// Check if the current COO element belongs to this row

if (row_index_coo[i] == row) {

dot += data_coo[i] * x[col_index_coo[i]];

}

}

// -----

// 3. Write the result back

// -----

y[row] += dot; }

}

data_coo

*	*	1	*
---	---	---	---

row_index_coo

*	*	2	*
---	---	---	---

col_index_coo

*	*	3	*
---	---	---	---



GPU Teaching Kit
Accelerated Computing



西南石油大学 计算机科学学院
SCHOOL OF COMPUTER SCIENCE, SOUTHWEST PETROLEUM UNIVERSITY




Module 10 – Parallel patterns: sparse matrix computation

10.5 Sorting and Partitioning for Regularization

Sorting and Partitioning for Regularization

- We can further reduce the filling cost **by sorting and partitioning the rows of sparse matrices**. This idea is to sort rows based on their length.
- The format is often referred to as the Jagged Diagonal Storage (**JDS**) format.


JDS



Sorting rows according to their length.

Benefit of JDS

- Transpose each section independently and launch a separate kernel on each section.
- Do not even need to launch a kernel for the section of rows with no nonzero elements.



JDS format and sectioned ELL


Row 0	3	0	1	0
Row 1	0	0	0	0
Row 2	0	2	4	1
Row 3	1	0	0	1

Nonzero values data[7] { 2, 4, 1, 3, 1, 1, 1 }

Column indices col_index[7] { 1, 2, 3, 0, 2, 0, 3 }

JDS row indices Jds_row_index[4] { 2, 0, 3, 1 }

Section pointers Jds_section_ptr[4] { 0, 3, 7, 7 }



JDS format and sectioned ELL.

SpMV kernel using JDS format

// spmv_jds: Sparse Matrix-Vector Multiplication (SpMV) using JDS format

```
__global__ void spmv_jds (  
    float * data,      // Non-zero values in JDS format  
    int * col_index,   // Column indices in JDS format  
    int * row_perm,    // Row permutation array(maps rows to original order)  
    int * row_nnz,     // The number of non-zero elements per row  
    float * x,         // Dense input vector x  
    float * y,         // Output vector y  
    int num_rows      // Total number of rows in the matrix  
) {  
    // Compute the row index this thread is responsible for  
    int row = blockDim.x * blockIdx.x + threadIdx.x;
```

SpMV kernel using JDS format

// spmv_jds: Sparse Matrix-Vector Multiplication (SpMV) using JDS format

...

```
if (row < num_rows) { // Ensure the thread processes a valid row
    float dot = 0.0f;
    int row_start = 0;
    // Compute the starting index for this row
    for (int j = 0; j < row; j++) {
        row_start += row_nnz[j];
    }
    // Compute the end index for non-zero elements in this row
    int row_end = row_start + row_nnz[row];
    for (int j = row_start; j < row_end; j++) {
        dot += data[j] * x[col_index[j]];
    }
    y[row_perm[row]] = dot;
}
```




GPU Teaching Kit

Accelerated Computing



西南石油大学 计算机科学学院

SCHOOL OF COMPUTER SCIENCE, SOUTHWEST PETROLEUM UNIVERSITY



The GPU Teaching Kit is licensed by NVIDIA and the University of Illinois under the [Creative Commons Attribution-NonCommercial 4.0 International License](https://creativecommons.org/licenses/by-nc/4.0/).